

Deep learning – Assignment 1

Image classification on CIFAR-100

Yonatan Zax - 204662779

We will start by naming some of the conclusion we have while writing the report.

Conclusions:

- **“Table of contents”** is a must while using any kind of notebook, might take time to write all the text cells, but the easy navigation is worth it at the end.
- **Data augmentation** was much more helpful than we expected at first, we will definitely use it in the future.
- **Printing some data samples** to inspect the augmentation changes and the resize changes were very helpful for understanding the benefit of using it.
- **Callbacks** are great time savers, by using EarlyStopping we were able to train our models with less unnecessary epochs and therefor spend less time waiting for the model’s fitting process.
- **Neptune.ai** is a great way to manage our Experiments. At some point it’s quite difficult to remember all the setting of your model, which one had better results. Neptune makes it easier to keep up with those experiments.

Table of Contents

0.Preparations:.....	2
1. Data exploration.....	2
2. Building a DL model.....	4
3. Trained model architecture.....	9

0. Preparations:


- Imports – from keras, sklearn and (pandas, numpy, matplotlib)
- Setting the Neptune account
- Load data – The data is available via keras.datasets

1. Data exploration

a. *Data size:* Train set has 50,000 records; Test set has 10,000 records.
Number of unique targets is 100

b. *What data does each sample contains:*

Each sample contains a single image and its classification.

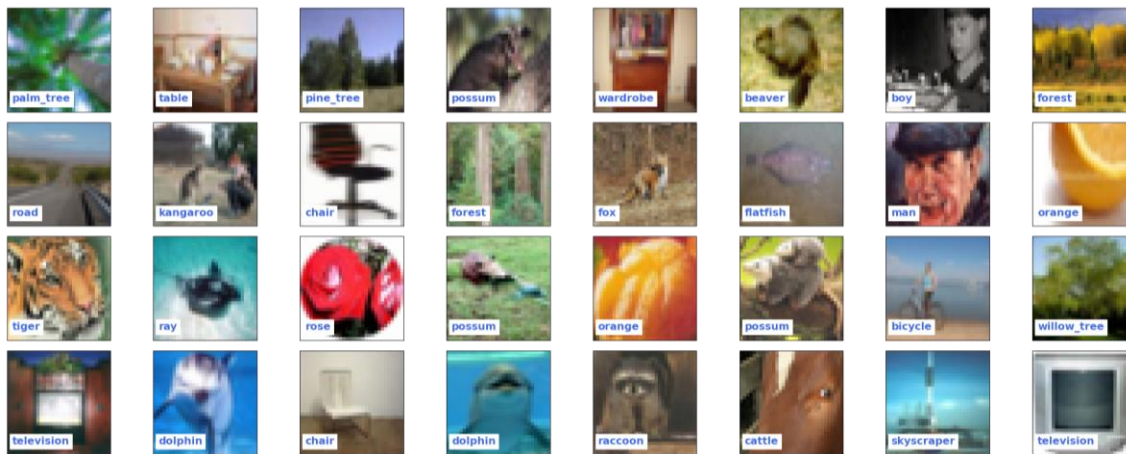


Dimensions
(32x32) pixels

Channels
RGB (3)

Num of classes
100

More Examples:



Preprocessing for the data? No, the data from keras.dataset is ready to use.

Augmentation: We can use augmentation, but the images are well cropped and the main element is centered therefore there is no need for augmentation for this dataset.

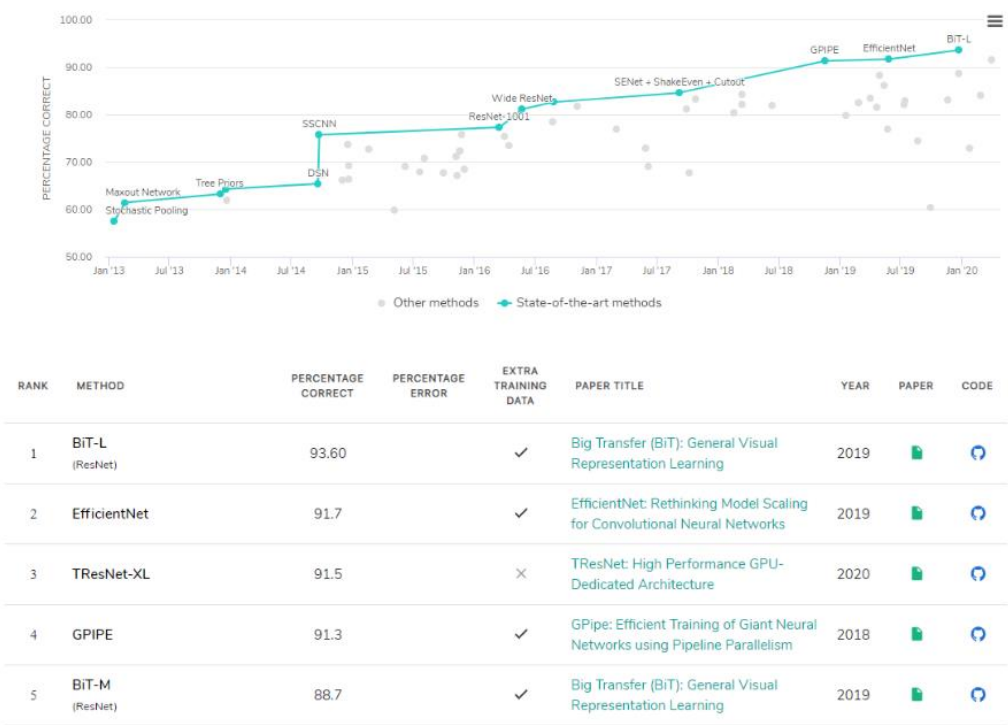
****Edit** – Later found out that using augmentation is also helpful and made a great usage as an improvement for the model.

c. *Is the data balanced?* Yes, the data is perfectly balanced - Each target has the exact same number of records (500 for the train, 100 for the test)

d. *Benchmark results on this dataset:*

[Link to Benchmarks on CIFAR-100](#)

Image Classification on CIFAR-100



We can see that the top 4 methods were able to get an accuracy score above 90%

Note that most of them were using 'Extra training data', this was not common in previous methods (from #6 downwards)

e. *Samples for each label (present easy and hard separable examples)*

Hard separable might be caused by similar backgrounds and/or colors of elements in the pictures, as well as similar shapes of the objects. Easy separable can be seen as images with different structures (colors, texture, backgrounds)



2. Building a DL model

- a. Validation strategy:* We used validation_split of 0.2
Train size will be 80% (40,000 records), Test size is 20% (10,000 records)
- b. Build model from scratch:*

Model's architecture:


Our model contains 3 blocks of Conv2D + Dropout + Max\Avg pooling + BatchNormalization. All Conv2D layers are using 'relu' as activation function. The first Conv2D has 16 filters, the last has 64. The last layer is a Dense layer of 100 because we have 100 different classes.

Layer (type)	Output Shape	Param #
Conv_16_3 (Conv2D)	(None, 30, 30, 16)	448
Dropout_0.1 (Dropout)	(None, 30, 30, 16)	0
MaxPooling (MaxPooling2D)	(None, 15, 15, 16)	0
BatchNormalization1 (BatchNo	(None, 15, 15, 16)	64
Conv_32_3 (Conv2D)	(None, 13, 13, 32)	4640
Dropout_0.2 (Dropout)	(None, 13, 13, 32)	0
Conv_64_3 (Conv2D)	(None, 11, 11, 64)	18496
Dropout_0.25 (Dropout)	(None, 11, 11, 64)	0
AvgPooling (AveragePooling2D	(None, 5, 5, 64)	0
BatchNormalization2 (BatchNo	(None, 5, 5, 64)	256
Flatten (Flatten)	(None, 1600)	0
Dense_100 (Dense)	(None, 100)	160100
Total params: 184,004		
Trainable params: 183,844		
Non-trainable params: 160		
Train on 40000 samples, validate on 10000 samples		

Training metrics:

----- Training score
Accuracy = 0.50
Loss = 1.86

----- Validation score
Validation Accuracy = 0.36
Validation Loss = 2.68

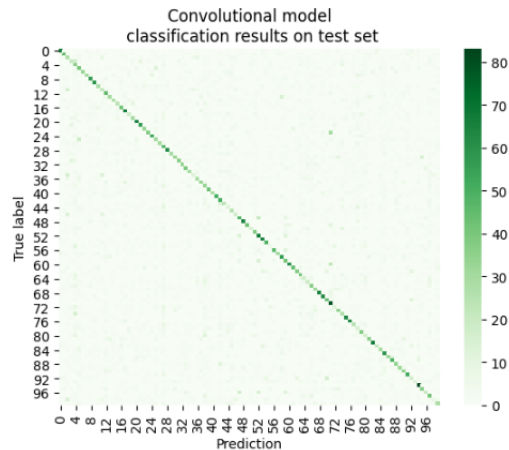


Evaluate the results:

There is a big difference between the train and validation caused by overfitting. As we can see, the Accuracy on the Train set is getting better each epoch while the validation is stabilized around 0.35. Same for the Loss, Train set is improving while the validation doesn't.

Basic model's prediction: Model's accuracy on test set is 38.7%

Confusion matrix: High values are located only on the diagonal line, meaning that there are no alternative classifications. Example: If there was high prediction outside of the diagonal line, it's a sign that the model is classifying one of the label as a different one. In our case, the model misses predictions over all targets.



Analyzing wrong predictions:

Some of the predictions were close



Other were far from being right



Even humans might find it difficult to classify the first row. Trees look alike and it's hard to get the correct one (same for Boy-Man and Squirrel-Rabbit)
The second row doesn't have any elements that might be similar between expected and actual prediction.

Comparing test set accuracy with validation accuracy

Validation
36.0 %

Test set
38.7 %

c. How to improve?

1. Build Conv2D layers up to size 256 (instead of 64) – Get a more complex model.
2. Use more epochs (instead of 10) – Let the model train more
3. Normalize the data, using `X_train.astype('float32')/255`

At first, we will show improves 1-3, later 4-5 will be added as well on another model.

4. Use Augmentation – Get different a variety of samples from a single image.
5. Add more data to training - Changing the split-validate ratio

d. Build an Improved model

Model's architecture:

Our model contains 5 blocks of Conv2D + Dropout + Max\Avg pooling + BatchNormalization. All Conv2D layers are using 'relu' as activation function. The first Conv2D has 16 filters, the last has 256. The last layer is a Dense layer of 100 because we have 100 different classes.

Layer (type)	Output Shape	Param #
Conv_16_3 (Conv2D)	(None, 30, 30, 16)	448
Dropout_0.2 (Dropout)	(None, 30, 30, 16)	0
MaxPooling (MaxPooling2D)	(None, 15, 15, 16)	0
BatchNormalization1 (BatchNo	(None, 15, 15, 16)	64
Conv_32_3 (Conv2D)	(None, 13, 13, 32)	4640
Dropout_0.15 (Dropout)	(None, 13, 13, 32)	0
BatchNormalization2 (BatchNo	(None, 13, 13, 32)	128
Conv_64_3 (Conv2D)	(None, 11, 11, 64)	18496
Dropout_0.25 (Dropout)	(None, 11, 11, 64)	0
BatchNormalization3 (BatchNo	(None, 11, 11, 64)	256
Conv_128_3 (Conv2D)	(None, 9, 9, 128)	73856
Dropout_0.3 (Dropout)	(None, 9, 9, 128)	0
AvgPooling (AveragePooling2D	(None, 4, 4, 128)	0
BatchNormalization4 (BatchNo	(None, 4, 4, 128)	512
Conv_256_3 (Conv2D)	(None, 2, 2, 256)	295168
Dropout_0.45 (Dropout)	(None, 2, 2, 256)	0
BatchNormalization5 (BatchNo	(None, 2, 2, 256)	1024
Flatten (Flatten)	(None, 1024)	0
Dense_100 (Dense)	(None, 100)	102500
Total params: 497,092		
Trainable params: 496,100		
Non-trainable params: 992		

Training metrics:

Due to the improves mentioned above, we got better results. Previous validation accuracy was 0.36. But, as seen in the graph below, the model is still overfitting.



Basic model's prediction: Model's accuracy on test set is 43.93%
Result on testset is exactly as the validation accuracy.

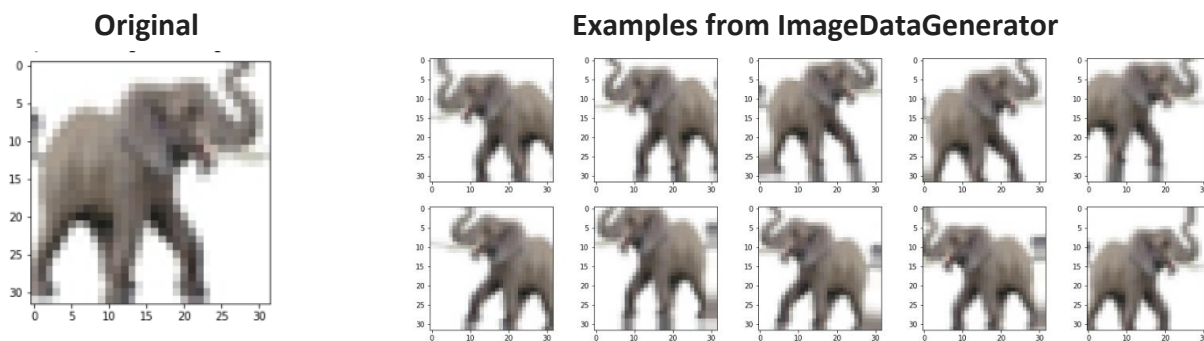
Improved model Version₂ (Using augmentation):

In order to overcome the overfitting, we will try to add more samples to the training data, did this by setting the validation_split to 0.9

In addition, we add ImageDataGenerator for augmentation usage. This way every image will be generated to a variety on samples at each epoch.


Augmentation parameters

1. rotation_range = 10 (Image will rotate at most 10 deg)
2. width_shift_range = 0.1 (Image will shift horizontally)
3. height_shift_range = 0.1 (Image will shift vertically)
4. shear_range = 0.01 (Shear angle in counter-clockwise direction in degrees)
5. zoom_range = 0.05 (Image will be zoomed-in)
6. horizontal_flip = True (Allowing horizontal flip)
7. vertical_flip = False (Vertical flip is not allowed)



Training metrics:

Due to more training samples we can see that both training and validation are more correlated.
Also, the validation accuracy is higher than before.




Basic model's prediction: Model's accuracy on test set is 48.76%
Result on testset is exactly as the validation accuracy.

3. Trained model architecture

We used the VGG-19 model with “imagenet” weights. This model is well known and was presented in our last lecturer.


At first the model wasn’t doing so well, after consulting with a friend who used this model before we were advised to drop the last 6 layers of the VGG-19 model.



a. Change the last layer:


We’ve added a flatten layer followed by a Dense output layer with 100 neurons. We used 20 epochs for fitting.

Evaluate the results



As we can see, the model is over fitting.

Confusion matrix: Accuracy on test set is 51.2%



d. Use the trained model as a feature extraction:

We used Random Forest Classifier, Accuracy on test set is 40%

e. Preform model stacking to improve performance

Improves

- 1. Increase the image size (Due to RAM limit, we will increase to 96x96)
- 2. Add layers to the model
 - o Add Dropout layer to handle overfitting
 - o Add BatchNormalization layers
- 3. Fit using more epochs

Resize image example:




Added layers on top of the VGG-19 architecture.

Flatten (Flatten)	(None, 73728)	0
Dense_512 (Dense)	(None, 512)	37749248
BatchNormalization (BatchNor	(None, 512)	2048
Dropout_0.3 (Dropout)	(None, 512)	0
Dense_256 (Dense)	(None, 256)	131328
Dropout_0.4 (Dropout)	(None, 256)	0
dense_7 (Dense)	(None, 100)	25700
=====		
Total params: 48,493,476		
Trainable params: 37,907,300		
Non-trainable params: 10,586,176		

By using EarlyStopping, the model fitting was over after 23/50 epochs.

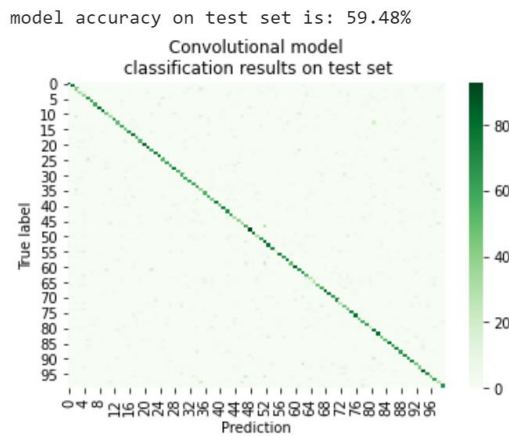
-----	Training score	-----
Accuracy = 0.61		
Loss = 1.37		
-----	Validation score	-----
Validation Accuracy = 0.60		
Validation Loss = 1.51		

Evaluate the results:



The model has good fit, validation score is very similar to the training. Moreover, the Accuracy we got (60%) is better than the model from scratch from section 2 (48%).

Confusion matrix



The accuracy on the test set is similar to the validation accuracy (60%)

Use as feature extractor:

We will use Random Forest Classifier, Accuracy on test set is 57%.
A significant improve than the previous model (40%)

*** Reminder – Conclusions are shown in the first page ***

Finally, hoping that you got this far, this assignment was indeed very long. We've learned a lot about image classification, approaches to improve accuracy and manage overfitting. We feel more confident about our understanding of the pipeline, from data exploration through preprocessing on to training a model and examining the results.